

For assignment 2 I have referenced assignment 1<sup>st</sup> “Blog” project.

Exercise 1:

## Schema Design and Indexing:

```
┆ aruzhan
class User(AbstractUser):
    bio = models.TextField(blank=True, null=True)

    ┆ aruzhan
    def __str__(self):
        return self.username
┆ aruzhan
class Tag(models.Model):
    name = models.CharField(max_length=50, unique=True)

    ┆ aruzhan
    def __str__(self):
        return self.name
```

```
┆ aruzhan
class Comment(models.Model):
    post = models.ForeignKey(Post, on_delete=models.CASCADE, related_name='comments')
    author = models.ForeignKey(User, on_delete=models.CASCADE)
    content = models.TextField(max_length=1000)
    created_date = models.DateTimeField(auto_now_add=True)

    ┆ aruzhan
    class Meta:
        indexes = [
            models.Index(fields=['post', 'created_date']),
        ]
    ┆ aruzhan
    def __str__(self):
        return self.content
```

```
┆ aruzhan *
class Post(models.Model):
    title = models.CharField(max_length=200)
    content = models.TextField()
    author = models.ForeignKey(User, on_delete=models.CASCADE)
    created_date = models.DateTimeField(auto_now_add=True)
    tags = models.ManyToManyField(to='Tag', related_name='posts')

    ┆ aruzhan *
    class Meta:
        indexes = [
            models.Index(fields=['author']),
            models.Index(fields=['tags']),
        ]
    ┆ aruzhan
    def __str__(self):
        return self.title
```

## Query Optimization:

- `Post.objects.prefetch_related('comments').all()` was used to request posts with all their comments in `post_list_view`
- The use of `select_related` might not be suitable because it only works with `ForeignKey` relationships. `Prefetch_related` is more suitable in this case.

## Optimization report:

-Here are views without `prefetch_related` command and the time that sql query took is 0.73ms.

```
@cache_page(5)
def post_list_view(request):
    posts = Post.objects.all().order_by('-created_date')
    paginator = Paginator(posts, per_page=5)
    page_number = request.GET.get('page')
    page_obj = paginator.get_page(page_number)
    return render(request, template_name='blog/post_list.html', context={'page_obj': page_obj})
```

```
def post_detail_view(request, id):
    post = get_object_or_404(Post, id=id)
    comments = post.comments.all().order_by('-created_date')
    comment_count = get_comment_count(post.id)
    if request.method == 'POST':
        form = CommentForm(request.POST)
        if form.is_valid():
            comment = form.save(commit=False)
            comment.author = request.user
            comment.post = post
            comment.save()
            return redirect(to='post_detail', id=post.id)
    else:
        form = CommentForm()
```

default 1.52 ms (6 queries including 2 similar and 2 duplicates)

Query	Timeline	Time (ms)	Action
SELECT *** FROM "blog_app_post" WHERE "blog_app_post"."id" = 14 LIMIT 21		0.73	<a href="#">Sel</a> <a href="#">Expl</a>
SELECT COUNT(*) AS "_count" FROM "blog_app_comment" WHERE "blog_app_comment"."post_id" = 14		0.14	<a href="#">Sel</a> <a href="#">Expl</a>
SELECT *** FROM "blog_app_comment" WHERE "blog_app_comment"."post_id" = 14 ORDER BY "blog_app_comment"."created_date" DESC		0.16	<a href="#">Sel</a> <a href="#">Expl</a>
SELECT *** FROM "blog_app_user" WHERE "blog_app_user"."id" = 4 LIMIT 21		0.18	<a href="#">Sel</a> <a href="#">Expl</a>
SELECT *** FROM "django_session" WHERE ("django_session"."expire_date" > "2024-10-10 11:21:47.510712" AND "django_session"."session_key" = "6hkk77zyatasjyxbwefuqta6itrag") LIMIT 21		0.22	<a href="#">Sel</a> <a href="#">Expl</a>
SELECT *** FROM "blog_app_user" WHERE "blog_app_user"."id" = 4 LIMIT 21		0.09	<a href="#">Sel</a> <a href="#">Expl</a>

2 similar queries. Duplicated 2 times.

2 similar queries. Duplicated 2 times.

2 similar queries. Duplicated 2 times.

Toggle Theme

History

/post\_detail/14/

Versions

Django 5.1.1

Time

CPU: 124.04ms (125.09ms)

Settings

Headers

Request

-Here is an example with `prefetch_related` ORM query used, and the time sql request took is 0.48ms

```

1 usage  ➤ aruzhan *
@cache_page(5)
def post_list_view(request):
    posts = Post.objects.prefetch_related('comments').all().order_by('-created_date')
    paginator = Paginator(posts, per_page=5)
    page_number = request.GET.get('page')
    page_obj = paginator.get_page(page_number)
    return render(request, template_name: 'blog/post_list.html', context: {'page_obj': page_obj})

def post_detail_view(request, id):
    post = get_object_or_404(Post.objects.prefetch_related('comments'), id=id)
    comments = post.comments.all().order_by('-created_date')
    comment_count = get_comment_count(post.id)
    if request.method == 'POST':
        form = CommentForm(request.POST)
        if form.is_valid():
            comment = form.save(commit=False)
            comment.author = request.user
            comment.post = post
            comment.save()
            return redirect(to: 'post_detail', id=post.id)
    else:
        form = CommentForm()

```

Recommendations for optimization:

- Using comment counting denormalization method:

```

1 usage  ➤ aruzhan
def get_comment_count(post_id):
    cache_key = f'post_{post_id}_comment_count'
    count = cache.get(cache_key)
    if count is None:
        count = Comment.objects.filter(post_id=post_id).count()
        cache.set(cache_key, count, timeout=60)
    return count

```

## Exercise 2:

### Basic Caching:

```
1 usage  ➤ aruzhan *
@cache_page(60)
def post_list_view(request):
    posts = Post.objects.prefetch_related('comments').all().order_by('-created_date')
    paginator = Paginator(posts, per_page=5)
    page_number = request.GET.get('page')
    page_obj = paginator.get_page(page_number)
    return render(request, template_name='blog/post_list.html', context={'page_obj': page_obj})
```

### Template Fragment Caching:

In post\_detail.html implemented caching of recent comments with timeout in 60 sec.

```
{% cache 60 recent_comments post.id %}
<h2>Recent Comments {{comment_count}}</h2>
{% for comment in comments %}
    <p>{{ comment.content }} by {{ comment.author.username }} on {{ comment.created_date }}</p>
{% endfor %}
{% endcache %}
```

### Low-level caching:

```
1 usage  ➤ aruzhan
def get_comment_count(post_id):
    cache_key = f'post_{post_id}_comment_count'
    count = cache.get(cache_key)
    if count is None:
        count = Comment.objects.filter(post_id=post_id).count()
        cache.set(cache_key, count, timeout=60)
    return count
```

### Caching backend:

Redis is implemented and used.

```
CACHES = {
    'default': {
        'BACKEND': 'django_redis.cache.RedisCache',
        'LOCATION': 'redis://127.0.0.1:6379/1',
        'OPTIONS': {
            'CLIENT_CLASS': 'django_redis.client.DefaultClient',
        }
    }
}
```

- For **measuring performance dependence on caching** I used Apache Benchmarking.

Without caching implementation:

```
- Time taken for tests: 45.818 seconds
- 100% 630 (longest request)
```

With caching:

```
- Time taken for tests: 45.294 seconds
- 100% 629 (longest request)
```

Conclusion: There is insignificant change in time because the database is not large enough.

Exercise 3:

**Nginx configuration:**

```
http{
    upstream django_servers {
        ip_hash;
        server 127.0.0.1:8000 max_fails=3 fail_timeout=30s;
        server 127.0.0.1:8001 max_fails=3 fail_timeout=30s;
        server 127.0.0.1:8002 backup;
    }
    server {
        listen 80;
        location /static/ {
            alias /home/sazanova/highload/assignment2/blog/staticfiles;
        }
        location / {
            proxy_pass http://127.0.0.1:8000;
            proxy_set_header Host $host;
            proxy_set_header X-Real-IP $remote_addr;
            proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
            proxy_set_header X-Forwarded-Proto $scheme;
        }
        location /health {
            proxy_pass http://django_servers/health;
            proxy_set_header Host $host;
            proxy_set_header X-Real-IP $remote_addr;
            proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
            proxy_set_header X-Forwarded-Proto $scheme;
        }
    }
}
```

I have created also servers for 8001 and 8002 ports.

## Round-robin load balancing:

I have implemented it in a new middleware.py file.

```
from django.http import HttpResponse

class RoundRobinMiddleware:
    def __init__(self, get_response):
        self.get_response = get_response

    def __call__(self, request):
        # Process the request (before the view is called)
        response = self.get_response(request)

        # Ensure response is not None
        if response is None:
            return HttpResponse(content="Empty Response", status=500)

        return response
```

Then added it into the MIDDLEWARE in settings.py:

```
'blog.middleware.RoundRobinMiddleware',
```

To enable **sticky sessions**, I used ip\_hash in nginx configuration file.

```
upstream django_servers {
    ip_hash;
    server 127.0.0.1:8000 max_fails=3 fail_timeout=30s;
    server 127.0.0.1:8001 max_fails=3 fail_timeout=30s;
    server 127.0.0.1:8002 backup;
}
```

**Health check** was implemented through writing limits for fails and time into “upstream” segment and writing a response function in view.py then adding it to the urls.py.

```
usage: aruzhan
def health_check(request):
    return JsonResponse({'status': 'ok'})

path('health/', views.health_check, name='health_check')
```

### Analysis of effectiveness of load balancer:

- It works properly and if any of the ports are killed or gone down it distributes tasks between ports.
- One user's actions are rightly distributed between ports and is transferred properly.

### Challenges faced:

- 1) There were some challenges in creating one static file directory for all css js files needed. I tried to implement django-debug-toolbar, but its css didn't work when running the project through gunicorn, but worked properly when running it through runserver. I have used the command `python manage.py collectstatic` and added all needed css, js files manually to `STATICFILES_DIRS` in settings.py
- 2) Problems dynamical elements in caching the page. When I cache the `post_list` page which is cached for 60 seconds, it doesn't instantly change login logout icons. I may solve it through implementing template fragment caching, but I needed to show how to use view level caching through `@cache_page(60)`.